

Course Overview

Introduction to Computer Systems

Overview

- Course theme
- Five realities
- How the course fits into the CS curriculum
- Academic integrity

Course Theme:

Abstraction Is Good But Don't Forget Reality

- **Most CS courses emphasize abstraction**
 - Abstract data types
 - Asymptotic analysis
- **These abstractions have limits**
 - Especially in the presence of bugs
 - Need to understand details of underlying implementations
- **Useful outcomes from taking this course**
 - Become more effective programmers
 - Able to find and eliminate bugs efficiently
 - Able to understand and tune for program performance
 - Prepare for later “systems” classes in CS
 - Computer Organization, Compilers, Operating Systems, Networks, Computer Architecture, Embedded Systems, Storage Systems, etc.

Great Reality #1:


Ints are not Integers, Floats are not Reals

■ Example 1: Is $x^2 \geq 0$?

■ Float's: Yes!

■ Int's:

- $40000 * 40000 \rightarrow 1600000000$
- $50000 * 50000 \rightarrow ??$



■ Example 2: Is $(x + y) + z = x + (y + z)$?

■ Unsigned & Signed Int's: Yes!

■ Float's:

- $(1e20 + -1e20) + 3.14 \rightarrow 3.14$
- $1e20 + (-1e20 + 3.14) \rightarrow ??$

Computer Arithmetic

■ Does not generate random values

- Arithmetic operations have important mathematical properties

■ Cannot assume all “usual” mathematical properties

- Due to finiteness of representations
- Integer operations satisfy “ring” properties
 - Commutativity, associativity, distributivity
- Floating point operations satisfy “ordering” properties
 - Monotonicity, values of signs

■ Observation

- Need to understand which abstractions apply in which contexts
- Important issues for compiler writers and serious application programmers

Great Reality #2:

You've Got to Know Assembly

- **Chances are, you'll never write programs in assembly**
 - Compilers are much better & more patient than you are
- **But: Understanding assembly is key to machine-level execution model**
 - Behavior of programs in presence of bugs
 - High-level language models break down
 - Tuning program performance
 - Understand optimizations done / not done by the compiler
 - Understanding sources of program inefficiency
 - Implementing system software
 - Compiler has machine code as target
 - Operating systems must manage process state
 - Creating / fighting malware
 - x86 assembly is the language of choice!

Great Reality #3: Memory Matters

Random Access Memory Is an Unphysical Abstraction

- **Memory is not unbounded**

- It must be allocated and managed
- Many applications are memory dominated

- **Memory referencing bugs especially pernicious**

- Effects are distant in both time and space

Memory Referencing Bug Example

```
typedef struct {
    int a[2];
    double d;
} struct_t;

double fun(int i) {
    volatile struct_t s;
    s.d = 3.14;
    s.a[i] = 1073741824; /* Possibly out of bounds */
    return s.d;
}
```

fun(0)	→	3.14
fun(1)	→	3.14
fun(2)	→	3.1399998664856
fun(3)	→	2.00000061035156
fun(4)	→	3.14
fun(6)	→	Segmentation fault

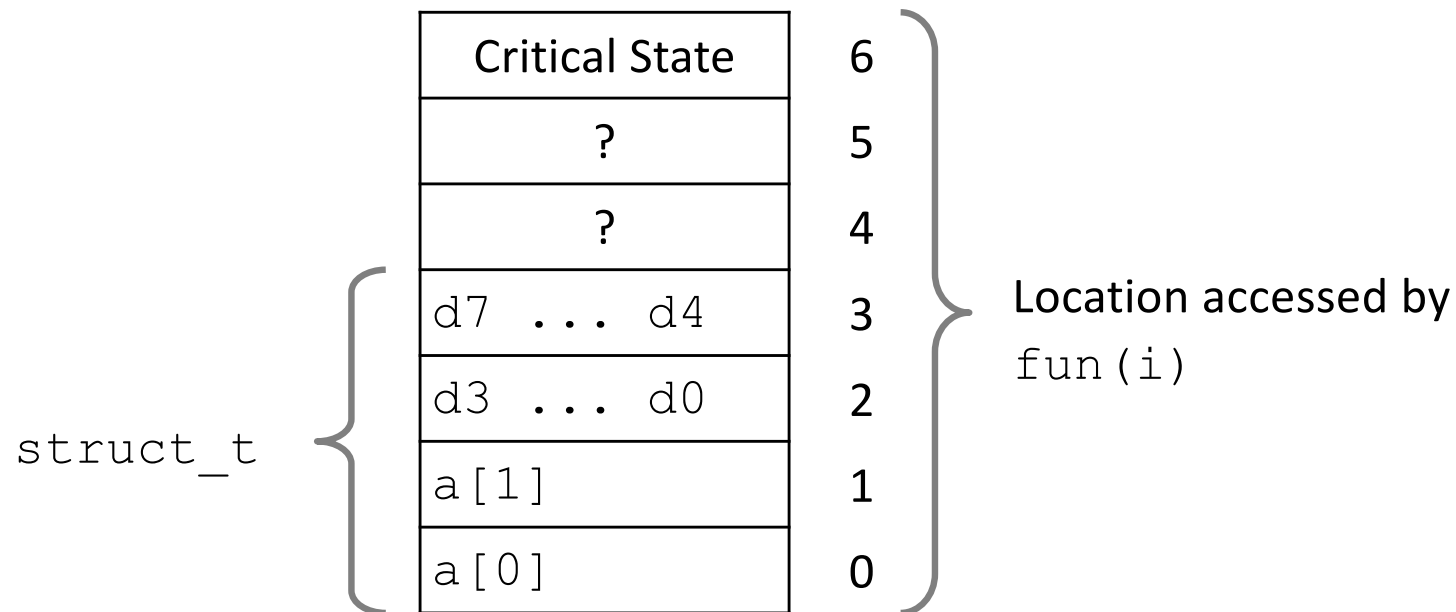
- Result is system specific

Memory Referencing Bug Example

```
typedef struct {  
    int a[2];  
    double d;  
} struct_t;
```

fun(0)	→	3.14
fun(1)	→	3.14
fun(2)	→	3.1399998664856
fun(3)	→	2.00000061035156
fun(4)	→	3.14
fun(6)	→	Segmentation fault

Explanation:



Memory Referencing Errors

■ C and C++ do not provide any memory protection

- Out of bounds array references
- Invalid pointer values
- Abuses of malloc/free

■ Can lead to nasty bugs

- Whether or not bug has any effect depends on system and compiler
- Action at a distance
 - Corrupted object logically unrelated to one being accessed
 - Effect of bug may be first observed long after it is generated

■ How can I deal with this?

- Program in Java, Python, Matlab, ...
- Understand what possible interactions may occur
- Use or develop tools to detect referencing errors (e.g. Valgrind)

Great Reality #3: Memory Matters

Random Access Memory Is an Unphysical Abstraction

- **Memory is not unbounded**

- It must be allocated and managed
- Many applications are memory dominated

- **Memory referencing bugs especially pernicious**

- Effects are distant in both time and space

- **Memory performance is not uniform**

- Cache and virtual memory effects can greatly affect program performance
- Adapting program to characteristics of memory system can lead to major speed improvements

Great Reality #4: There's more to performance than asymptotic complexity

- **Constant factors matter too!**
- **And even exact op count does not predict performance**
 - Easily see 10:1 performance range depending on how code written
 - Must optimize at multiple levels: algorithm, data representations, procedures, and loops
- **Must understand system to optimize performance**
 - How programs compiled and executed
 - How to measure program performance and identify bottlenecks
 - How to improve performance without destroying code modularity and generality

Memory System Performance Example

```
void copyij(int src[2048][2048],
            int dst[2048][2048])
{
    int i,j;
    for (i = 0; i < 2048; i++)
        for (j = 0; j < 2048; j++)
            dst[i][j] = src[i][j];
}
```

4.3ms


```
void copyji(int src[2048][2048],
            int dst[2048][2048])
{
    int i,j;
    for (j = 0; j < 2048; j++)
        for (i = 0; i < 2048; i++)
            dst[i][j] = src[i][j];
}
```

81.8ms

2.0 GHz Intel Core i7 Haswell

- Hierarchical memory organization
- Performance depends on access patterns
 - Including how step through multi-dimensional array

Why The Performance Differs



Great Reality #5:

Computers do more than execute programs

- **They need to get data in and out**
 - I/O system critical to program reliability and performance
- **They communicate with each other over networks**
 - Many system-level issues arise in presence of network
 - Concurrent operations by autonomous processes
 - Coping with unreliable media
 - Cross platform compatibility
 - Complex performance issues

Course Perspective

- Most Systems Courses are Builder-Centric
 - Computer Organization and Architecture
 - Design pipelined processor in Verilog
 - Operating Systems
 - Implement sample portions of operating system
 - Compilers
 - Write compiler for simple language
 - Networking
 - Implement and simulate network protocols

Course Perspective (Cont.)

■ Our Course is Programmer-Centric

- Purpose is to show that by knowing more about the underlying system, one can be more effective as a programmer
- Enable you to
 - Write programs that are more reliable and efficient
 - Incorporate features that require hooks into OS
 - E.g., concurrency, signal handlers
- Cover material in this course that you won't see elsewhere
- Not just a course for dedicated hackers
 - **We bring out the hidden hacker in everyone!**

Cheating: Description

■ What is cheating?

- Sharing code: by copying, retyping, **looking at**, or supplying a file
- Describing: verbal description of code from one person to another.
- Coaching: helping your friend to write a lab, line by line
- Searching the Web for solutions
- Copying code from a previous course or online solution
 - You are only allowed to use code we supply

■ What is NOT cheating?


- Explaining how to use systems or tools
- Helping others with high-level design issues

Textbooks

- Randal E. Bryant and David R. O'Hallaron,
 - *Computer Systems: A Programmer's Perspective, Third Edition* (CS:APP3e), Pearson, 2016
 - This book really matters for the course!
 - How to solve labs
 - Practice problems typical of exam problems



- Brian Kernighan and Dennis Ritchie,
 - *The C Programming Language, Second Edition*, Prentice Hall, 1988
 - Still the best book about C, from the originators



Course Components

■ Lectures

- Higher level concepts

■ Labs (4-5)

- The heart of the course
- 1-3 weeks each
- Provide in-depth understanding of an aspect of systems
- Programming and measurement

■ Flipped Classroom

■ Final Exam

- Test your understanding of concepts & mathematical principles

往届同学课程学习心得

可以让人对计算机体系有了一个总体的了解，这门课程不仅让后续课程的学习变得轻松了许多(其他班的同学学体系结构时看见汇编都被吓了大跳)，也让人学习了很多实际中有用的东西。比如说实验就要求安装Linux来使用。

感觉对具体实现有个初步概念，然后一步一步很有成就感，解开谜团的感觉。对将来的电脑学习充满了兴趣。

计算机系统基础是让我开始觉得计算机真的有意思的课程。做labs虽然耗费精力但是好玩。对计算机组成原理的过渡效果非常好，计组学习基本没有太大压力。尤其是优化程序性能的那一部分，对编程帮助也蛮大的。

别整虚的（比如相信书光看就能看懂，大佬除外），多动手做课后题和lab

计算机系统基础是CMU的镇校神课当然要上



Other Rules of the Lecture Hall

- Laptops: permitted
- Electronic communications: **forbidden**
 - No email, instant messaging, cell phone calls, etc
- Presence in lectures: voluntary, recommended by me, while requested by the school

Assistance

- 通知、答疑：QQ群/私信
- 作业与实验平台：计算机学院希冀平台
 - <https://cslabcg.whu.edu.cn/>
 - 账号已导入，**务必登陆后补充email**，方便找回密码
- 教辅TA：QQ群，批改作业，答疑
- 线下：教师办公室时间
 - 每周四中午12点~2点

Policies: Grading

- Exams (60%)
- Labs (20%): weighted according to effort
- Classroom presentation (10%)
- Homework (10%)

Labs

- L0 (clab): Basic C programming skills
- L1 (datalab): Manipulating bits
- L2 (bomblab): Defusing a binary bomb
- L3 (attacklab): The basics of code injection attacks
- L4 (shelllab): A simple Unix shell program with job control
- L5 (proxylab): A concurrent caching Web proxy that sits between the browser and the rest of the World Wide Web

如何学好一门系统课程？



如何学好一门系统课程？



- 在陌生的环境中醒来→想知道自己在哪里？
- 用卷尺测量桌子的高度（91cm），一支试管（高密度、可以忽略空气阻力），用秒表记录它掉落的时间（0.37s），重复实验20次，得到平均值0.348s。 $d=1/2gt^2$ ，所以 $g=15m/s^2$ ，是 $9.8m/s^2$ 的1.5倍→我不在地球上！
- 我可能在一台离心机里？
- 一卷尼龙绳，卷尺，做一个单摆，无论幅度多大，摆动一个来回的时间——周期——是恒定的，只取决于单摆的长度和重力加速度。高处，346周，10分钟；低处，346周，10分钟！如果在离心机里，你离中心越远，向心力越大，结果没有差别，至少没有达到影响单摆的周期。
- 开始排除超大离心机的可能性（半径至少700m，转速88m/s，会有涡流、风噪）
- 在另一颗行星？太阳系内没有任何一颗行星、卫星或小行星有这么大引力，除非在木星的风暴里，否则不会体验到这么大的重力加速度....

Welcome
&
Enjoy!